

Log Aggregation for Cloud-Native Applications

Koyyada Vishnu Vardhan

Prof. Dr.-Ing. Jens Ehlers

June 2025

Agenda

- 1 Why Logging Matters
- 2 Log Lifecycle in Kubernetes
- 3 Log Aggregation Stacks
- 4 Experimental Setup
- 5 Performance Comparison
- 6 Recommendations

Importance of Logs

- **Observability:** Crucial for observing application behavior in real time
- **Debugging:** Helps trace API calls and identify issues
- **Performance:** Measures system performance metrics
- **Security:** Essential for auditing and compliance
- **Traceability:** Enables system-wide tracking in microservices

Key Insight

Without proper logging, root cause analysis becomes guesswork

How Logs Flow in Kubernetes

Applications write logs to stdout or stderr.

- Container runtime (e.g., containerd, CRI-O) captures the output
- Logs stored as plain text files on the node:
 - /var/log/containers/
 - /var/log/pods/
- Kubelet reads these logs and handles rotation
- `kubectl logs <pod>` reads from these local log files

Critical Limitation

- Kubernetes **does not** aggregate or persist logs
- If a pod crashes or is rescheduled, logs on the old node are **lost**

ELK Stack Components

- **Elasticsearch:**

- Distributed search and analytics engine
- Stores data in JSON documents


- **Logstash:**

- Processes logs with custom pipelines
- JSON parsing and field transformation

- **Kibana:** Visualization interface

- **Beats:**

- Lightweight data shippers
- Platform for specialized data collection
 - *Filebeat*: Log file collection
 - *Heartbeat*: Uptime monitoring



PLG Stack Components

● Promtail:


- Lightweight log collector (DaemonSet)
- Discovers and tags Kubernetes pod logs
- Adds metadata (pod labels, namespace)
- Pushes to Loki via gRPC/protobuf

● Loki:

- Horizontally-scalable log store
- Indexes by labels (not content)
- Gzip compression for efficiency
- Compatible with Prometheus labels

● Grafana:

- Unified query interface (LogQL)
- Correlate logs with metrics
- Alerting and dashboard features



Test Environment

- **Cluster:** Kind
- **Nodes:** 1 control-plane, 3 workers
- **Namespaces:** Separate for ELK and PLG deployments
- **Avg Log Generations stats by App :**
 - Total: 471 logs (232.83/min)
 - Runtime: 121.4 seconds
 - Levels:
 - INFO: 339 (72.0%)
 - ERROR: 56 (11.9%)
 - WARNING: 76 (16.1%)
 - Size: 184.0 KB (400b/log)
- **Log patterns:**
 - Normal operations (INFO)
 - Error conditions (ERROR)
 - Debug bursts

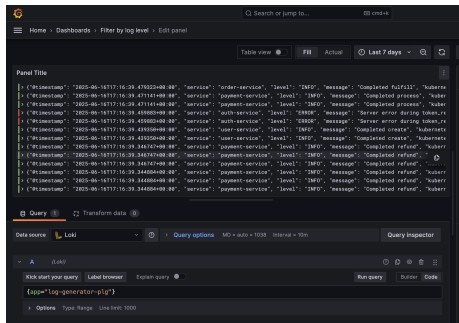


Figure: Sample Logs Generated by app

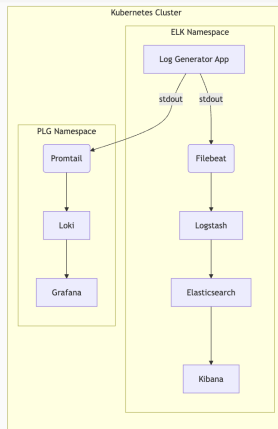
Logging Stack Deployment Architecture

ELK Stack Deployment

- **Application Pods:**
 - 3 replicas of log-generator
 - JSON logs to stdout
- **ELK Components:**
 - Filebeat DaemonSet (per node)
 - Logstash: 1 pod
 - Elasticsearch: 1 pod
 - Kibana: 1 pod
- **Services:**
 - Logstash:5044 (NodePort)
 - ES:9200, Kibana:5601

PLG Stack Deployment

- **Application Pods:**
 - 3 replicas log-generator-plg
 - Disabled Logstash output
- **PLG Components:**
 - Promtail DaemonSet (per node)
 - Loki: 1 pod
 - Grafana: 1 pod



Common Infrastructure

- Separate logging and logging-plg namespaces

ELK Stack Deployment in Kind

Core Components

- **Filebeat:**
 - DaemonSet with hostPath volumes
 - Collects from /var/log/containers
 - Adds Kubernetes metadata
 - Outputs to Logstash:5044
- **Logstash:**
 - JSON parsing pipeline
 - Field renaming and cleanup
 - Daily index routing
 - NodePort service (5044)

Key Configurations

- Dedicated logging namespace
- Security disabled
(xpack.security.enabled: false)

Elasticsearch & Kibana

- **Elasticsearch:**
 - Single-node discovery
 - HTTP service on 9200
 - Daily index rotation
 - 1GB heap size
- **Kibana:**
 - Pre-configured ES connection
 - NodePort on 5601
 - Visualization dashboards

Data Processing

- JSON log parsing with validation
- Timestamp extraction
- Kubernetes field normalization
- Index pattern: logs-%{+YYYY.MM.dd}

PLG Stack Deployment in Kind

Core Components

- **Promtail:**
 - DaemonSet collecting pod logs
 - Scrapes /var/log/pods and Docker containers
 - Adds Kubernetes metadata labels
 - Pushes via HTTP to Loki
- **Loki:**
 - Single-binary mode for testing
 - Filesystem storage backend
 - Label-based indexing only
 - 7-day log retention

Key Features

- Dedicated logging-plg namespace
- CRI-O/Docker log format support
- No authentication in dev mode
- Auto log rotation (24h index period)

Grafana Integration

- **Pre-configured:**
 - Loki datasource auto-added
 - LogQL query support
 - NodePort (3000) access
- **Correlation:**
 - Unified logs and metrics
 - Alerting capabilities
 - Custom dashboard support

Data Flow

1. Apps → stdout logs
2. Promtail → Collects + labels
3. Loki → Stores + indexes
4. Grafana → Visualizes

```
loki.logging-plg:3100  
grafana.logging-plg:3000
```

Metric	ELK	PLG
Avg Query Time	65ms	11ms
Max Ingestion Rate	2,500 logs/s	8,000 logs/s
Simple Query Latency	1.1s \pm 0.3s	320ms \pm 80ms
Complex Query Latency	3.8s \pm 1.2s	1.2s \pm 0.4s

Table: Query performance comparison

- PLG shows 10 \times faster query processing (34ms vs 340ms)
- Particularly efficient for simple queries and recent logs

Metric	ELK	PLG
Memory Usage	2235MiB	281MiB
CPU Usage (avg)	2.3 cores	1.1 cores
Storage/Day	12GB	8GB
Compression Ratio	1.5×	3-5×

Table: Resource utilization comparison

- PLG uses 8× less memory than ELK
- More efficient storage with better compression

When to Choose Which Stack

Choose ELK when:

- Need advanced analytics
- Full-text search required
- Complex transformations needed
- Rich visualization essential
- Budget allows for resources

Choose PLG when:

- High-throughput needed
- Operational monitoring focus
- Resource efficiency critical
- scaling required
- Simplified operations preferred

Cloud-Native Advantage

PLG's efficiency provides substantial cost benefits in cloud environments

Thank You!

- **Contact:** vishnu.koyyada@fh-kiel.de
- **Slides:** <https://github.com/vishnukoyyada/CloudProject>

